

For columns:

*Glucose,*  
*BloodPressure,*  
*BMI,*  
*DiabetesPedigreeFunction*

If the column value is 0, then they should be considered as **missing data**.

1. Firstly, replace all Missing values with relevant figures.

```
import numpy as np
import pandas as pd
df = pd.read_csv('Dataset_Day7.csv')
print(df.info())
missing_value_percent = df.isna().sum() / len(df)
* 100
print(missing_value_percent)
skewness = df.skew()
print(skewness)
df["Glucose"].fillna(df["Glucose"].median(),
inplace=True)
df["BloodPressure"].fillna(df["BloodPressure"].me
an(), inplace=True)
df["BMI"].fillna(df["BMI"].median(),
inplace=True)
df["Outcome"].fillna(df["Outcome"].mean(),
inplace=True)
print(df.info())
```

```

1 C:\Users\tejas\PycharmProjects\pythonProject\venv\Scripts\python.exe C:\Users\
  tejas\PycharmProjects\pythonProject\START\Day8Q1.py
2 <class 'pandas.core.frame.DataFrame'>
3 RangeIndex: 768 entries, 0 to 767
4 Data columns (total 7 columns):
5 #   Column                Non-Null Count  Dtype
6 ---  ---
7 0   Pregnancies            768 non-null    int64
8 1   Glucose                768 non-null    int64
9 2   BloodPressure          768 non-null    int64
10 3   BMI                   768 non-null    float64
11 4   DiabetesPedigreeFunction 768 non-null    float64
12 5   Age                   768 non-null    int64
13 6   Outcome                768 non-null    int64
14 dtypes: float64(2), int64(5)
15 memory usage: 42.1 KB
16 None
17 Pregnancies            0.0
18 Glucose                0.0
19 BloodPressure          0.0
20 BMI                   0.0
21 DiabetesPedigreeFunction 0.0
22 Age                   0.0
23 Outcome                0.0
24 dtype: float64
25 Pregnancies            0.901674
26 Glucose                0.173754
27 BloodPressure          -1.843608
28 BMI                   -0.428982
29 DiabetesPedigreeFunction 1.919911
30 Age                   1.129597
31 Outcome                0.635017
32 dtype: float64
33 <class 'pandas.core.frame.DataFrame'>
34 RangeIndex: 768 entries, 0 to 767
35 Data columns (total 7 columns):
36 #   Column                Non-Null Count  Dtype
37 ---  ---
38 0   Pregnancies            768 non-null    int64
39 1   Glucose                768 non-null    int64
40 2   BloodPressure          768 non-null    int64
41 3   BMI                   768 non-null    float64
42 4   DiabetesPedigreeFunction 768 non-null    float64
43 5   Age                   768 non-null    int64
44 6   Outcome                768 non-null    int64
45 dtypes: float64(2), int64(5)
46 memory usage: 42.1 KB
47 None
48
49 Process finished with exit code 0

```

2. Then remove all existing outliers and get the final data for classification.

```

import numpy as np
import pandas as pd
df = pd.read_csv('Dataset_Day7.csv')
print(df.info())
missing_value_percent = df.isna().sum() / len(df)

```

```

* 100
print(missing_value_percent)
skewness = df.skew()
print(skewness)
df["Glucose"].fillna(df["Glucose"].median(),
inplace=True)
df["BloodPressure"].fillna(df["BloodPressure"].me
an(), inplace=True)
df["BMI"].fillna(df["BMI"].median(),
inplace=True)
df["Outcome"].fillna(df["Outcome"].mean(),
inplace=True)
print(df.info())
OutlierData = pd.DataFrame()
temp = df[["Pregnancies", "Glucose",
"BloodPressure", "BMI", "DiabetesPedigreeFunction"
, "Age", "Outcome" ]]
for col in ["Pregnancies", "Glucose",
"BloodPressure", "BMI", "DiabetesPedigreeFunction"
, "Age", "Outcome"]:
    Q1 = temp[col].quantile(0.25) # Gives 25th
Percentile or Q1
    Q3 = temp[col].quantile(0.75) # Gives 75th
Percentile or Q3

    IQR = Q3 - Q1

    UpperBound = Q3 + 1.5 * IQR
    LowerBound = Q1 - 1.5 * IQR

    OutlierData[col] = temp[col][(temp[col] <
LowerBound) | (temp[col] > UpperBound)]
    print(len(OutlierData))
    df_OutlierFree = df.drop(OutlierData.index,
axis=0)
    df_OutlierFree.info()

```

```

tejas\PycharmProjects\pythonProject\START\Day8Q2.py
2 <class 'pandas.core.frame.DataFrame'>
3 RangeIndex: 768 entries, 0 to 767
4 Data columns (total 7 columns):
5 #   Column                Non-Null Count  Dtype
6 ---  ---
7 0   Pregnancies            768 non-null    int64
8 1   Glucose                768 non-null    int64
9 2   BloodPressure          768 non-null    int64
10 3   BMI                   768 non-null    float64
11 4   DiabetesPedigreeFunction 768 non-null    float64
12 5   Age                   768 non-null    int64
13 6   Outcome               768 non-null    int64
14 dtypes: float64(2), int64(5)
15 memory usage: 42.1 KB
16 None
17 Pregnancies            0.0
18 Glucose                0.0
19 BloodPressure          0.0
20 BMI                   0.0
21 DiabetesPedigreeFunction 0.0
22 Age                   0.0
23 Outcome               0.0
24 dtype: float64
25 Pregnancies            0.901674
26 Glucose                0.173754
27 BloodPressure          -1.843608
28 BMI                   -0.428982
29 DiabetesPedigreeFunction 1.919911
30 Age                   1.129597
31 Outcome               0.635017
32 dtype: float64
33 <class 'pandas.core.frame.DataFrame'>
34 RangeIndex: 768 entries, 0 to 767
35 Data columns (total 7 columns):
36 #   Column                Non-Null Count  Dtype
37 ---  ---
38 0   Pregnancies            768 non-null    int64
39 1   Glucose                768 non-null    int64
40 2   BloodPressure          768 non-null    int64
41 3   BMI                   768 non-null    float64
42 4   DiabetesPedigreeFunction 768 non-null    float64
43 5   Age                   768 non-null    int64
44 6   Outcome               768 non-null    int64
45 dtypes: float64(2), int64(5)
46 memory usage: 42.1 KB
47 None
48 4
49 <class 'pandas.core.frame.DataFrame'>
50 Index: 764 entries, 0 to 767
51 Data columns (total 7 columns):
52 #   Column                Non-Null Count  Dtype
53 ---  ---
54 0   Pregnancies            764 non-null    int64
55 1   Glucose                764 non-null    int64
56 2   BloodPressure          764 non-null    int64
57 3   BMI                   764 non-null    float64
58 4   DiabetesPedigreeFunction 764 non-null    float64

```

```

59 5 Age 764 non-null int64
60 6 Outcome 764 non-null int64
61 dtypes: float64(2), int64(5)
62 memory usage: 47.8 KB
63 4
64 <class 'pandas.core.frame.DataFrame'>
65 Index: 764 entries, 0 to 767
66 Data columns (total 7 columns):
67 # Column Non-Null Count Dtype
68 ---
69 0 Pregnancies 764 non-null int64
70 1 Glucose 764 non-null int64
71 2 BloodPressure 764 non-null int64
72 3 BMI 764 non-null float64
73 4 DiabetesPedigreeFunction 764 non-null float64
74 5 Age 764 non-null int64
75 6 Outcome 764 non-null int64
76 dtypes: float64(2), int64(5)
77 memory usage: 47.8 KB
78 4
79 <class 'pandas.core.frame.DataFrame'>
80 Index: 764 entries, 0 to 767
81 Data columns (total 7 columns):
82 # Column Non-Null Count Dtype
83 ---
84 0 Pregnancies 764 non-null int64
85 1 Glucose 764 non-null int64
86 2 BloodPressure 764 non-null int64
87 3 BMI 764 non-null float64
88 4 DiabetesPedigreeFunction 764 non-null float64
89 5 Age 764 non-null int64
90 6 Outcome 764 non-null int64
91 dtypes: float64(2), int64(5)
92 memory usage: 47.8 KB
93 4
94 <class 'pandas.core.frame.DataFrame'>
95 Index: 764 entries, 0 to 767
96 Data columns (total 7 columns):
97 # Column Non-Null Count Dtype
98 ---
99 0 Pregnancies 764 non-null int64
100 1 Glucose 764 non-null int64
101 2 BloodPressure 764 non-null int64
102 3 BMI 764 non-null float64
103 4 DiabetesPedigreeFunction 764 non-null float64
104 5 Age 764 non-null int64
105 6 Outcome 764 non-null int64
106 dtypes: float64(2), int64(5)
107 memory usage: 47.8 KB
108 4
109 <class 'pandas.core.frame.DataFrame'>
110 Index: 764 entries, 0 to 767
111 Data columns (total 7 columns):
112 # Column Non-Null Count Dtype
113 ---
114 0 Pregnancies 764 non-null int64
115 1 Glucose 764 non-null int64
116 2 BloodPressure 764 non-null int64
117 3 BMI 764 non-null float64

```

```

118 4 DiabetesPedigreeFunction 764 non-null float64
119 5 Age 764 non-null int64
120 6 Outcome 764 non-null int64
121 dtypes: float64(2), int64(5)
122 memory usage: 47.8 KB
123 4
124 <class 'pandas.core.frame.DataFrame'>
125 Index: 764 entries, 0 to 767
126 Data columns (total 7 columns):
127 # Column Non-Null Count Dtype
128 --- ---
129 0 Pregnancies 764 non-null int64
130 1 Glucose 764 non-null int64
131 2 BloodPressure 764 non-null int64
132 3 BMI 764 non-null float64
133 4 DiabetesPedigreeFunction 764 non-null float64
134 5 Age 764 non-null int64
135 6 Outcome 764 non-null int64
136 dtypes: float64(2), int64(5)
137 memory usage: 47.8 KB
138 4
139 <class 'pandas.core.frame.DataFrame'>
140 Index: 764 entries, 0 to 767
141 Data columns (total 7 columns):
142 # Column Non-Null Count Dtype
143 --- ---
144 0 Pregnancies 764 non-null int64
145 1 Glucose 764 non-null int64
146 2 BloodPressure 764 non-null int64
147 3 BMI 764 non-null float64
148 4 DiabetesPedigreeFunction 764 non-null float64
149 5 Age 764 non-null int64
150 6 Outcome 764 non-null int64
151 dtypes: float64(2), int64(5)
152 memory usage: 47.8 KB
153
154 Process finished with exit code 0
155

```

3, Split the data into 70% training and 30% testing data. Then, use a k-Nearest Neighbor algorithm with target variable as 'Outcome'.

- Print the default model performance metrics: Accuracy, Precision, Recall, F1Score
- Plot a Precision & Recall vs k(no. of neighbours) curve (both Prec and Rec on the same graph). Find the k for which F1-score is the highest. **Use any one Distance Metric for this problem.**
- Find the best distance metric, no. of neighbors combination for the kNN algorithm

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```

from sklearn.model_selection import
train_test_split
from sklearn.neighbors import
KNeighborsClassifier
from sklearn.metrics import accuracy_score,
precision_score, recall_score, f1_score
import math

df = pd.read_csv('Dataset_Day7.csv')
print(df.info())
missing_value_percent = df.isna().sum() / len(df)
* 100
print(missing_value_percent)
skewness = df.skew()
print(skewness)
df["Glucose"].fillna(df["Glucose"].median(),
inplace=True)
df["BloodPressure"].fillna(df["BloodPressure"].me
an(), inplace=True)
df["BMI"].fillna(df["BMI"].median(),
inplace=True)
df["Outcome"].fillna(df["Outcome"].mean(),
inplace=True)
print(df.info())
OutlierData = pd.DataFrame()
temp = df[["Pregnancies", "Glucose",
"BloodPressure", "BMI",
"DiabetesPedigreeFunction", "Age", "Outcome"]]
for col in ["Pregnancies", "Glucose",
"BloodPressure", "BMI",
"DiabetesPedigreeFunction", "Age", "Outcome"]:
    Q1 = temp[col].quantile(0.25) # Gives 25th
Percentile or Q1
    Q3 = temp[col].quantile(0.75) # Gives 75th
Percentile or Q3

    IQR = Q3 - Q1

    UpperBound = Q3 + 1.5 * IQR
    LowerBound = Q1 - 1.5 * IQR

    OutlierData[col] = temp[col][(temp[col] <

```

```

LowerBound) | (temp[col] > UpperBound)]

df_OutlierFree = df.drop(OutlierData.index,
axis=0)
df_OutlierFree.info()

X = df_OutlierFree.drop('Outcome', axis=1) # all
columns except 'Outcome'
y = df_OutlierFree['Outcome'] # target Variable
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.3,
random_state=203)
unique_classes = np.unique(y)
print('Number of unique classes:',
len(unique_classes))
k_start = int(math.sqrt(len(X_train)))
print(k_start)
metric_start = 'euclidean'

knn = KNeighborsClassifier(n_neighbors=k_start,
metric=metric_start)
# fit the model
knn = knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

print("Model Performance metrics are as below:")
print("Accuracy: ", accuracy_score(y_test,
y_pred))
print("Precision: ", precision_score(y_test,
y_pred, zero_division=1))
print("Recall: ", recall_score(y_test, y_pred))
print("F1-Score: ", f1_score(y_test, y_pred))

k_values = np.arange(2,25)
metric_values = ['euclidean', 'manhattan',
'hamming']
prec = []
rec = []
acc = []
PerfData = pd.DataFrame(columns=['Nearest
Neighbor', 'Distance Metric', 'Precision',

```



```

'Recall', 'Accuracy'])

for dm in metric_values:
    for k in k_values:
        knn = KNeighborsClassifier(n_neighbors=k,
metric=dm)
        knn = knn.fit(X_train, y_train)
        y_pred = knn.predict(X_test)

        row = [[k, dm, precision_score(y_test,
y_pred, zero_division=1), recall_score(y_test,
y_pred),
                    accuracy_score(y_test, y_pred)]]

        df2 = pd.DataFrame(row, columns=['Nearest
Neighbor', 'Distance Metric', 'Precision',
'Recall', 'Accuracy'])
        PerfData = pd.concat([PerfData, df2],
ignore_index=True)

print(PerfData.tail())
precision = PerfData['Precision']
recall = PerfData['Recall']
PerfData["F1 Score"] = (2 * PerfData["Precision"]
* PerfData["Recall"]) / (PerfData["Precision"] +
PerfData["Recall"])
print(PerfData[PerfData['F1 Score'] ==
max(PerfData['F1 Score'])])
# Plot Precision & Recall vs k Curve
plt.plot(k_values, PerfData[PerfData["Distance
Metric"]=="euclidean"]['Precision'],
label='Precision')
plt.plot(k_values, PerfData[PerfData["Distance
Metric"]=="euclidean"]['Recall'], label='Recall')
plt.xlabel('k (Number of Neighbors)')
plt.ylabel('Precision / Recall')
plt.title('Precision & Recall vs k Curve')
plt.legend(loc='upper right')
plt.show()

```

```

2 <class 'pandas.core.frame.DataFrame'>
3 RangeIndex: 768 entries, 0 to 767
4 Data columns (total 7 columns):
5 #   Column                Non-Null Count  Dtype
6 ---  ---
7 0   Pregnancies            768 non-null    int64
8 1   Glucose                 768 non-null    int64
9 2   BloodPressure           768 non-null    int64
10 3   BMI                     768 non-null    float64
11 4   DiabetesPedigreeFunction 768 non-null    float64
12 5   Age                     768 non-null    int64
13 6   Outcome                 768 non-null    int64
14 dtypes: float64(2), int64(5)
15 memory usage: 42.1 KB
16 None
17 Pregnancies            0.0
18 Glucose                 0.0
19 BloodPressure           0.0
20 BMI                     0.0
21 DiabetesPedigreeFunction 0.0
22 Age                     0.0
23 Outcome                 0.0
24 dtype: float64
25 Pregnancies            0.901674
26 Glucose                 0.173754
27 BloodPressure          -1.843608
28 BMI                     -0.428982
29 DiabetesPedigreeFunction 1.919911
30 Age                     1.129597
31 Outcome                 0.635017
32 dtype: float64
33 <class 'pandas.core.frame.DataFrame'>
34 RangeIndex: 768 entries, 0 to 767
35 Data columns (total 7 columns):
36 #   Column                Non-Null Count  Dtype
37 ---  ---
38 0   Pregnancies            768 non-null    int64
39 1   Glucose                 768 non-null    int64
40 2   BloodPressure           768 non-null    int64
41 3   BMI                     768 non-null    float64
42 4   DiabetesPedigreeFunction 768 non-null    float64
43 5   Age                     768 non-null    int64
44 6   Outcome                 768 non-null    int64
45 dtypes: float64(2), int64(5)
46 memory usage: 42.1 KB
47 None
48 <class 'pandas.core.frame.DataFrame'>
49 Index: 764 entries, 0 to 767
50 Data columns (total 7 columns):
51 #   Column                Non-Null Count  Dtype
52 ---  ---
53 0   Pregnancies            764 non-null    int64
54 1   Glucose                 764 non-null    int64
55 2   BloodPressure           764 non-null    int64
56 3   BMI                     764 non-null    float64
57 4   DiabetesPedigreeFunction 764 non-null    float64
58 5   Age                     764 non-null    int64

```

File - Day8Q3

```
59 6 Outcome 764 non-null int64
60 dtypes: float64(2), int64(5)
61 memory usage: 47.8 KB
62 Number of unique classes: 2
63 23
64 Model Performance metrics are as below:
65 Accuracy: 0.7478260869565218
66 Precision: 0.7222222222222222
67 Recall: 0.47560975609756095
68 F1-Score: 0.5735294117647057
69 Nearest Neighbor Distance Metric Precision Recall Accuracy
70 64 20 hamming 1.0 0.0 0.643478
71 65 21 hamming 1.0 0.0 0.643478
72 66 22 hamming 1.0 0.0 0.643478
73 67 23 hamming 1.0 0.0 0.643478
74 68 24 hamming 1.0 0.0 0.643478
75 Nearest Neighbor Distance Metric Precision Recall Accuracy F1 Score
76 15 17 euclidean 0.754386 0.52439 0.769565 0.618705
77
78 Process finished with exit code 0
79
```

Figure 1

